# Company Overview

February 2025

Except for the historical information contained herein, certain matters in this presentation including, but not limited to, statements as to: expectations with respect to growth, performance and benefits of NVIDIA's products, services, and technologies, including Blackwell, and related trends and drivers; expectations with respect to supply and demand for NVIDIA's products, services, and technologies, including Blackwell, and related matters including inventory, production and distribution; our financial position; projected market growth and trends, market opportunity, demand, and growth drivers; our financial and business outlook; our dividend program; expectations with respect to NVIDIA's third party arrangements, including with its collaborators and partners; third parties adopting our products and technologies; expectations with respect to technology developments and related trends and drivers; expectations with respect to AI and related industries; our sustainability goals; and other statements that are not historical facts are forward-looking statements.

These forward-looking statements and any other forward-looking statements that go beyond historical facts that are made in this presentation are subject to risks and uncertainties that may cause actual results to differ materially. Important factors that could cause actual results to differ materially include: global economic and political conditions; NVIDIA's reliance on third parties to manufacture, assemble, package and test NVIDIA's products; the impact of technological development and competition; development of new products and technologies or enhancements to NVIDIA's existing product and technologies; market acceptance of NVIDIA's products or NVIDIA's partners' products; design, manufacturing or software defects; changes in consumer preferences and demands; changes in industry standards and interfaces; unexpected loss of performance of NVIDIA's products or technologies when integrated into systems and other factors; and changes in applicable laws and regulations.

NVIDIA has based these forward-looking statements largely on its current expectations and projections about future events and trends that it believes may affect its financial condition, results of operations, business strategy, short-term and long-term business operations and objectives, and financial needs. These forward-looking statements are subject to a number of risks and uncertainties, and you should not rely upon the forward-looking statements as predictions of future events. The future events and trends discussed in this presentation may not occur and actual results could differ materially and adversely from those anticipated or implied in the forward-looking statements. Although NVIDIA believes that the expectations reflected in the forward-looking statements are reasonable, the company cannot guarantee that future results, levels of activity, performance, achievements or events and circumstances reflected in the forward-looking statements will occur. Except as required by law, NVIDIA disclaims any obligation to update these forward-looking statements to reflect future events or circumstances. For a complete discussion of factors that could materially affect NVIDIA's financial results and operations, please refer to the reports we file from time to time with the SEC, including NVIDIA's most recent Annual Report on Form 10-K, Quarterly Reports on Form 10-Q, and Current Reports on Form 8-K. Copies of reports we file with the SEC are posted on NVIDIA's website and are available from NVIDIA without charge

Many of the products and features described herein remain in various stages and will be offered on a when-and-if-available basis. The statements within are not intended to be, and should not be interpreted as a commitment, promise, or legal obligation, and the development, release, and timing of any features or functionalities described for our products is subject to change and remains at the sole discretion of NVIDIA. NVIDIA will have no liability for failure to deliver or delay in the delivery of any of the products, features or functions set forth herein.

NVIDIA uses certain non-GAAP measures in this presentation including non-GAAP operating income, non-GAAP operating margin, and free cash flow. NVIDIA believes the presentation of its non-GAAP financial measures enhances investors' overall understanding of the company's historical financial performance. The presentation of the company's non-GAAP financial measures is not meant to be considered in isolation or as a substitute for the company's financial results prepared in accordance with GAAP, and the company's non-GAAP measures may be different from non-GAAP measures used by other companies. Further information relevant to the interpretation of non-GAAP financial measures, and reconciliations of these non-GAAP financial measures to the most comparable GAAP measures, may be found in the slide titled "Reconciliation of Non-GAAP to GAAP Financial Measures."



# NVIDIA

NVIDIA's invention of the GPU in 1999 sparked the growth of the PC gaming market, redefined computer graphics, revolutionized accelerated computing, ignited the era of modern AI, and is fueling industrial digitalization across markets.

Today, two transitions are occurring simultaneously—accelerated computing and generative AI—transforming the computer industry and every other industry worldwide, and NVIDIA is enabling these transitions with our full-stack computing platform and data-center-scale offerings.

NVIDIA's platform is installed in several hundred million computers, is available in every cloud and from every server maker, powers over 75% of the TOP500 supercomputers, and has ~5.9 million developers.

Headquarters: Santa Clara, CA | Headcount: ~36,000

# NVIDIA’s Accelerated Computing Platform

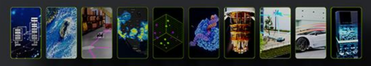
## Data center scale innovation across chips, networking, systems, software, and algorithms



NIM
CUDA-Accelerated
Agentic AI Libraries



Omniverse CUDA-Accelerated
Physical AI Libraries

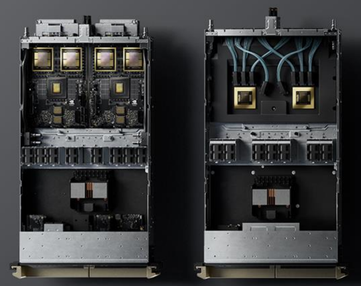


CUDA-X Libraries

CUDA • DOCA • NCCL
Cluster-Scale Software
System Software
Chip Software

Accelerated
Software Stack

GB200 NVL72 SuperPOD

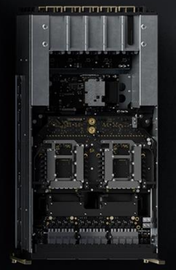


Grace Blackwell
MGX Node



NVLink Switch





Quantum Switch



Spectrum-X
Switch



Chips Purpose-Built for AI Supercomputing
GPU | CPU | DPU | NIC | NVLink Switch | IB Switch | ENET Switch

NVIDIA has accelerated software and compute by a 1,000,000X in the last decade, far surpassing Moore’s law.

Accelerated computing requires full-stack innovation—optimizing across every layer of computing—from chips and systems to software and algorithms, demanding deep understanding of the problem domain.

Our platform extends from the cloud and enterprise data centers to supercomputing, edge computing, PCs, and robotics.

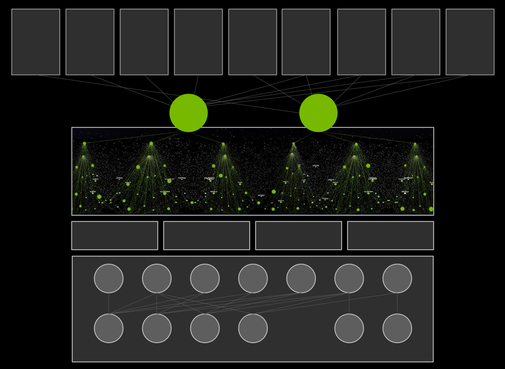
# What Is Accelerated Computing?

## A full-stack approach: silicon, systems, software

Not just a superfast chip—accelerated computing is a full-stack combination of:

- Chip(s) with specialized processors
- Algorithms in acceleration libraries
- Domain experts to refactor applications

To speed up compute-intensive parts of an application



**Amdahl's law:**

The overall system speed-up (S) gained by optimizing a single part of a system by a factor (s) is limited by the proportion of execution time of that part (p).

$$S = \frac{1}{(1-p) + \frac{p}{s}}$$

For example:

- If 90% of the runtime can be accelerated by 100X, the application is sped up 9X
- If 99% of the runtime can be accelerated by 100X, the application is sped up 50X
- If 80% of the runtime can be accelerated by 500X, or even 1,000X, the application is sped up 5X

# Why Accelerated Computing?

## Advancing computing in the post-Moore's law era

Accelerated computing is needed to tackle the most impactful opportunities of our time—like AI, climate simulation, drug discovery, ray tracing, and robotics.

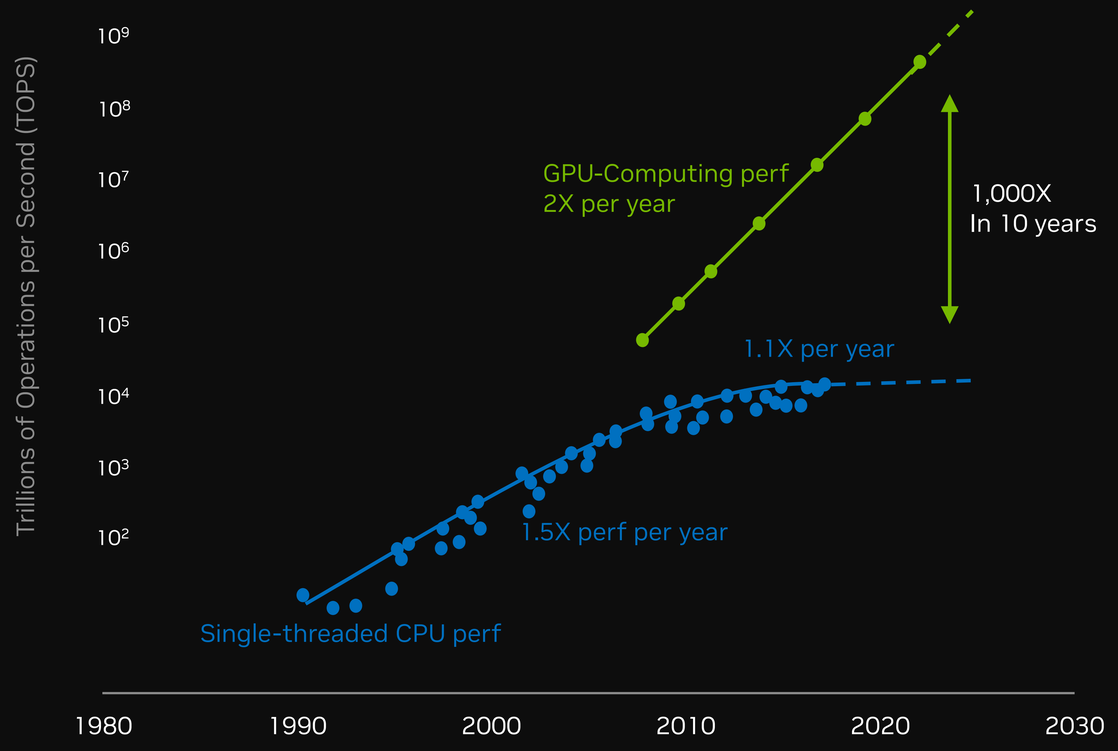
NVIDIA is uniquely dedicated to accelerated computing —working top-to-bottom, refactoring applications and creating new algorithms, and bottom-to-top inventing new specialized processors, like RT Cores and Tensor Cores.

*"It's the end of Moore's law as we know it."*

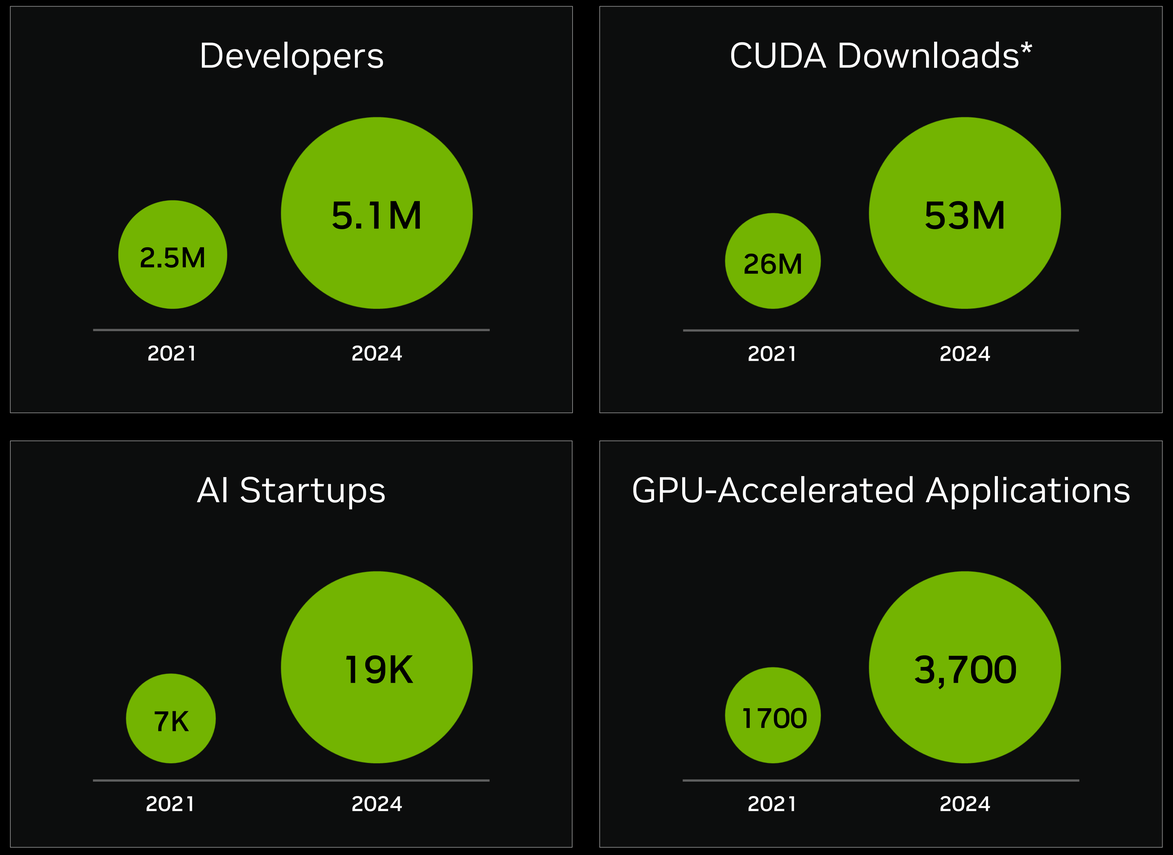
—John Hennessy, Oct 2018

*"Moore's law is dead."*

—Jensen Huang, GTC 2013



# NVIDIA’s Accelerated Computing Ecosystem




- The NVIDIA accelerated computing platform has attracted the largest ecosystem of developers, supporting a rapidly growing universe of applications and industry innovation.
- Developers can engage with NVIDIA through CUDA—our parallel computing programming model introduced in 2006—or at higher layers of the stack, including libraries, pretrained AI models, SDKs, and other development tools.

*Cumulative

# AI Driving a Powerful Investment Cycle and Significant Returns

## AI can augment creativity and productivity by orders of magnitude across industries

**AI Agents** will take action to automate tasks at superhuman speed, transforming businesses and freeing workers to focus on other tasks.

**Copilots** based on LLMs will generate documents, answer questions, or summarize missed meetings, emails, and chats—adding hours of productivity per week. Specialized for fields such as **software development, legal services** or **education** and can boost productivity by as much as 50%.

**Social media, search, and e-commerce apps** are using deep recommenders to offer more relevant content and ads to their customers, increasing engagement and monetization.

**Creators** can generate stunning, photorealistic images with a single text prompt—compressing workflows that take days or weeks into minutes in industries from advertising to game development.

**Call center agents** augmented with AI chatbots can dramatically increase productivity and customer satisfaction.

**Drug discovery and financial services** are seeing order-of-magnitude workflow acceleration from AI.

**Manufacturing** workflows are reinvented and automated through generative AI and robotics, boosting productivity.

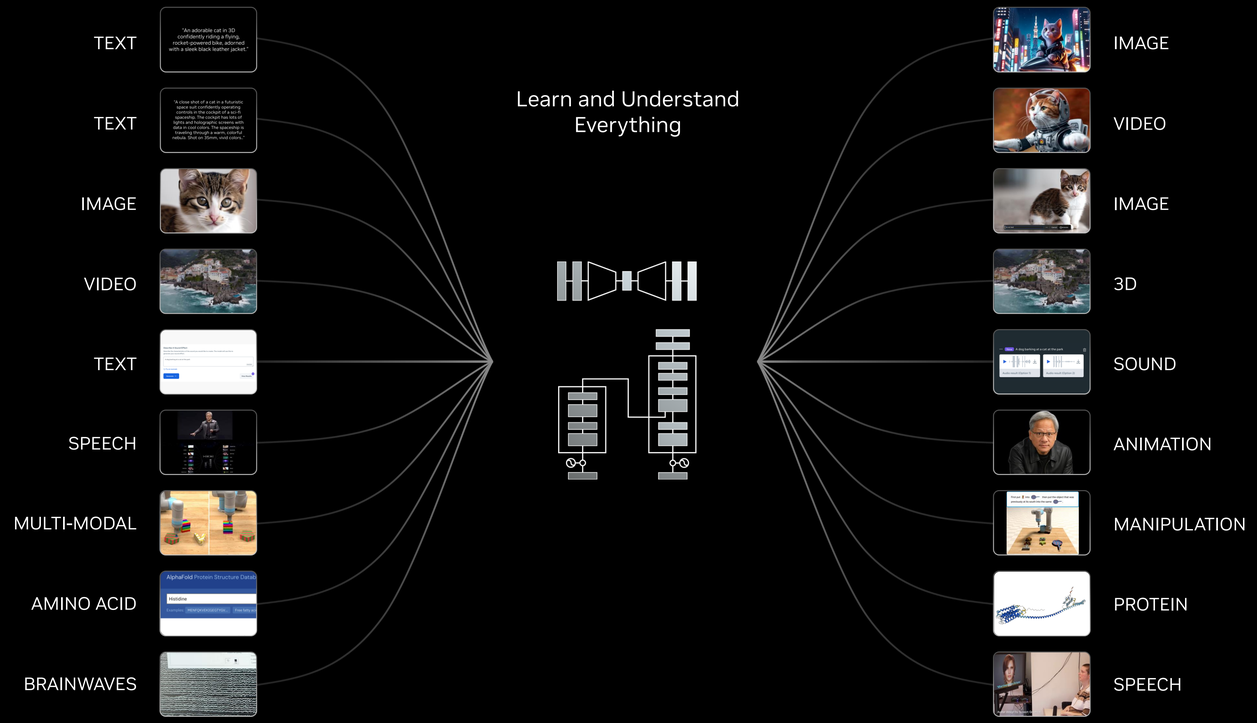



NVIDIA

*Source: Goldman Sachs, Cowen, Statista, Capital One, Wall Street Journal, Resource Watch, NVIDIA internal analysis*

# Generative AI

## The most important computing platform of our generation




The era of generative AI has arrived, unlocking new opportunities for AI across many different applications.

Generative AI is trained on large amounts of data to find patterns and relationships, learning the representation of almost anything with structure.

It can then be prompted to generate text, images, video, code, or even proteins.

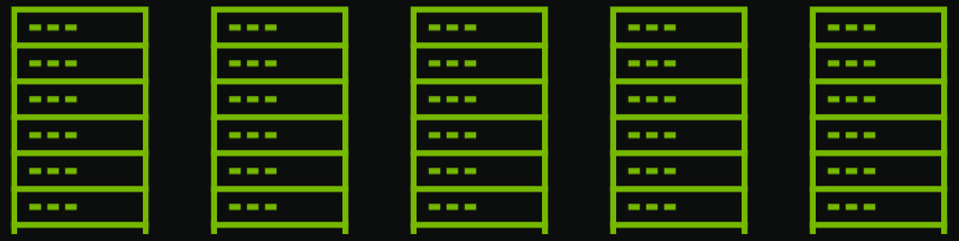
For the very first time, computers can augment the human ability to generate information and create.

1,600+ generative AI companies are building on NVIDIA.

# Delivering Tremendous Value to Customers

Significant reduction in TCO with each generation

Hopper | 8,000 GPUs | 15MW



Blackwell GB200 NVL72 | 2,000 GPUs | 4MW



**4X**

Reduction in Power

*LLM Training Workload: GPT-MoE-1.8T | Train in 90 days | H100 vs GB200 NVL72*

NVIDIA

# AI Factories—A New Class of Data Centers

## Production of digital intelligence tokens

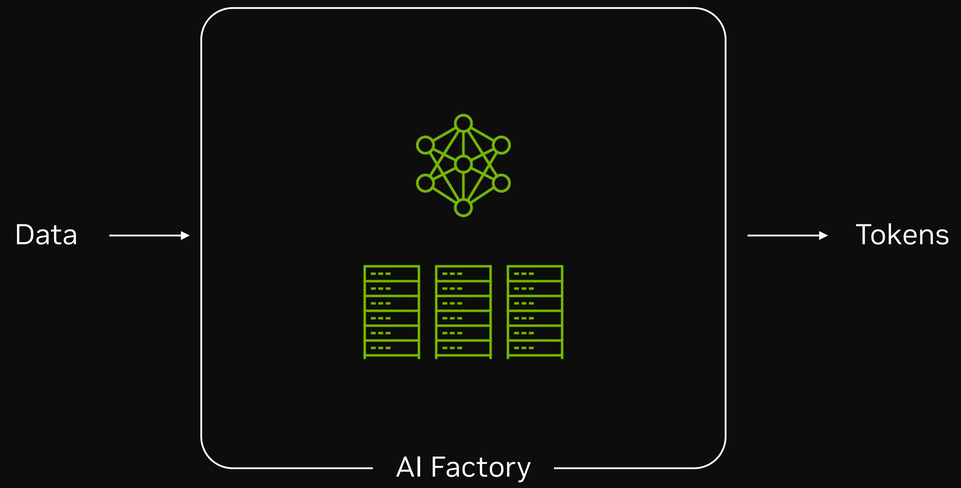
**AI factories** are a new form of computing infrastructure. Their purpose is not to store user and company data or run ERP and CRM applications. AI factories are highly optimized systems purpose-built to process raw data, refine it into models, and produce monetizable tokens with great scale and efficiency.

In the AI industrial revolution, data is the raw material, tokens are the new commodity, and NVIDIA is the token generator in the AI factory.

Every company will produce digital intelligence. Tokens will be transformed into intelligent responses and actions of digital nurses, tutors, customer service agents, chip designers, manufacturing robots and autonomous cars, Weather prediction agents will warn us of storms. Some companies will build and operate AI factories, while others will rent.

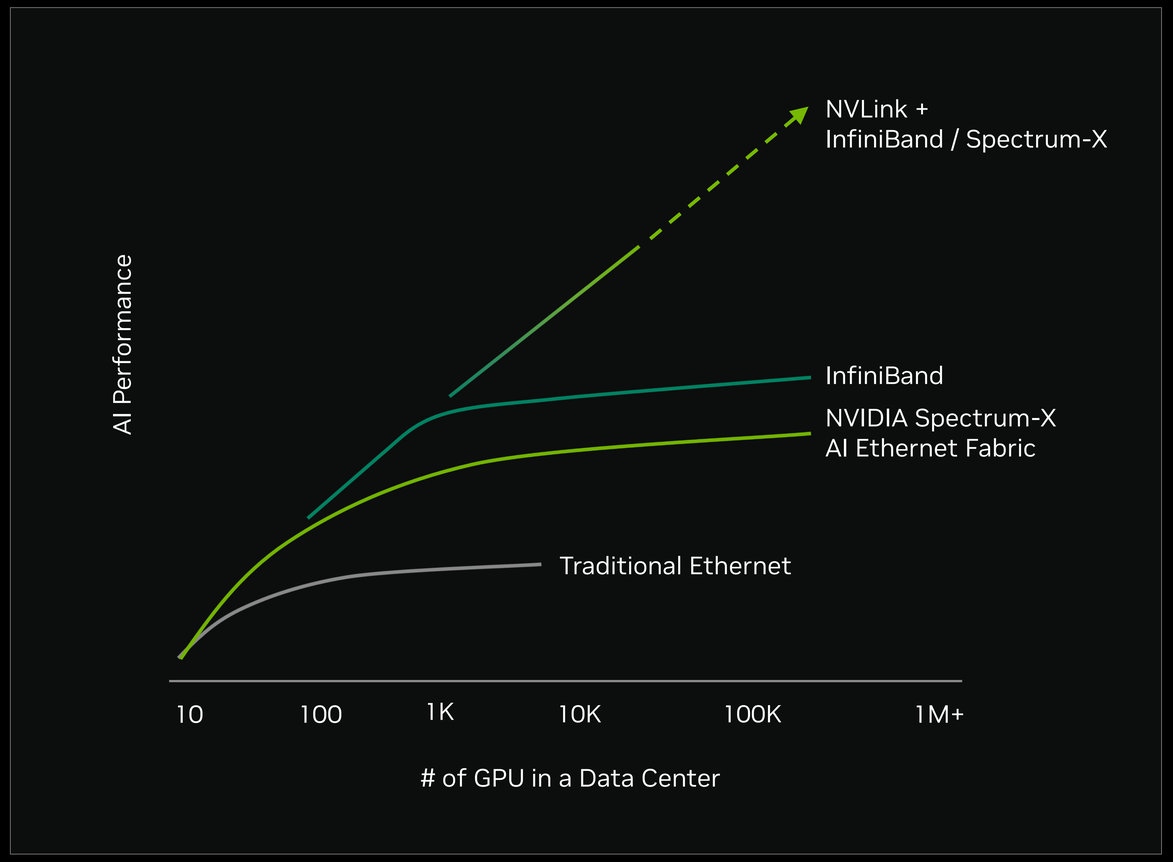
Countries are awakening to the need to treat their data as a national resource and AI factories as an essential national infrastructure. Data encodes a nation's history, knowledge, and culture, and can be transformed into the sovereign AI for its companies, startups, universities, and governments.

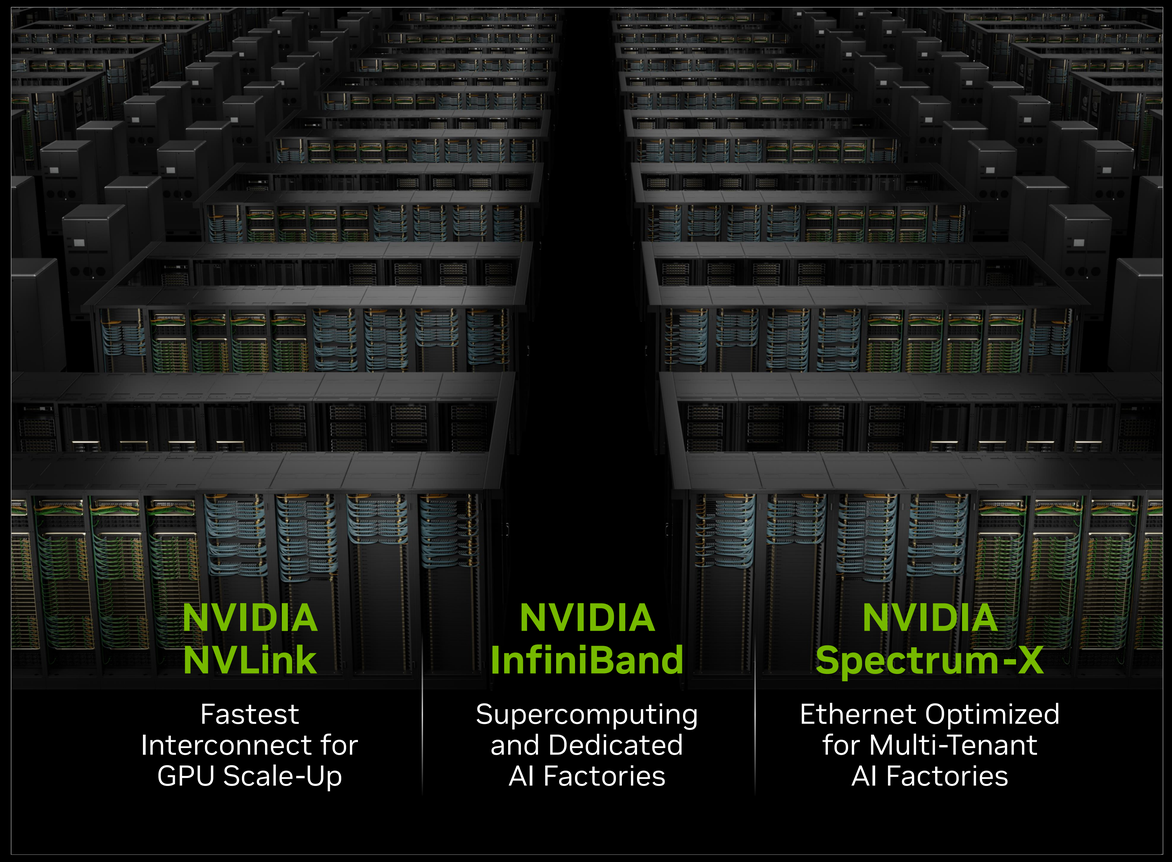
NVIDIA builds the complete AI system and licenses **NVIDIA AI Enterprise**, the AI stack and operating system for AI factories.



# Extending NVIDIA Networking to Scale Up & Scale Out AI in Any Data Center

New NVLink and Spectrum-X increase networking opportunity beyond InfiniBand to every data center




## Generative AI Is a Data Center-Scale Computing Workload

Limitless scaling with NVLINK + InfiniBand or Spectrum-X

# Accelerated Computing Starts With CUDA Libraries

## Delivering up to 200X speedup across major workloads

Unlike CPU general-purpose computing, GPU-accelerated computing requires software and algorithms to be redesigned. Software is not automatically accelerated in the presence of a GPU or accelerator.

NVIDIA CUDA libraries encapsulate NVIDIA-engineered algorithms that enable applications to be accelerated on NVIDIA's installed base. They deliver dramatically higher performance—compared to CPU-only alternatives—across application domains, including AI and high-performance computing, and significantly reduce runtime, cost, and energy, while increasing scale.

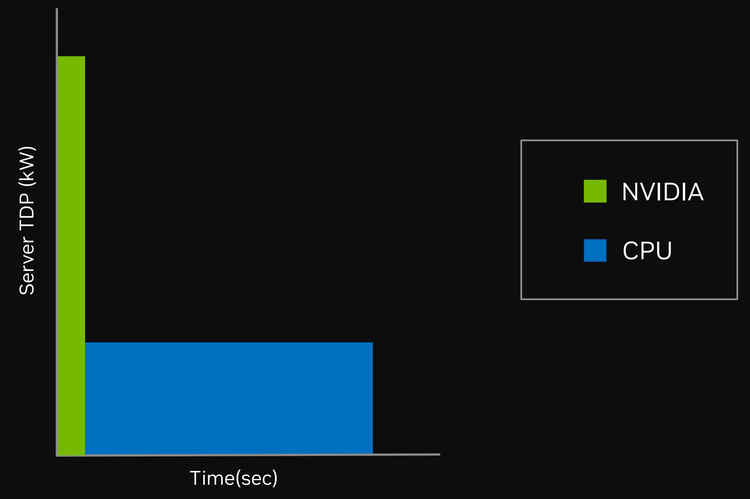
With over 400 CUDA libraries, NVIDIA can address many major workloads across a wide range of industries. As new libraries become available, they unlock new markets adding to our long-term opportunity.

| | | |
|---|---|---|
| ~200X | **Data Processing** | cuVS, cuDF-Spark, cuDF-pandas, cuDF-Polars, cuGraph, cuML, XGBoost, RAPIDS, NeMo Curator, cuSOLVER, cuIO |
| ~200X | **Computer Vision** | CV-CUDA, Deepstream, TAO, Holoscan, cuCIM, TensorRT, Triton Inference Server, DALI, nvImageCodec, cuDNN, nvJPEG, nvJPEG2000, nvTIFF, NPP, Video Codec SDK, Magnum IO, NCCL, cuVS, DALI |
| ~100X | **Science** | Earth-2 CorrDiff, Holoscan, Parabricks, Monai, Modulus, Warp, cuLitho, cuQuantum, CUDA-Q, AmgX, cuDSS, cuFFT, cuSOLVER, cuBLAS, cuSPARSE, cuTENSOR, cuGraph, Magnum IO, NCCL, NVSHMEM, RAFT, cuNumeric, Sionna |
| ~100X | **Deep Learning** | cuDNN, CUTLASS, Megatron, TensorRT, TRT LLM, NCCL, NV-Triton, CUDA-optimized PyTorch, Tensorflow, Triton, Jax |
| ~50X | **Recommender Systems** | Merlin, HugeCTR, TensorRT, Triton Inference Server, cuBLAS, cuDNN, cuFFT, cuSPARSE, CUTLASS, Magnum IO, NCCL, cuVS |
| ~30X | **Speech AI** | Riva, TensorRT, Triton Inference Server, NeMo, cuBLAS, cuDNN, cuFFT, CUTLASS |
| ~100X | **Agentic and Physical AI** | ACE, Riva, Nemo, Tokkio Digital Human, Holoscan, Metropolis, Omniverse, Isaac, DRIVE, cuLitho, cuMotion, cuOpt, Aerial CUDA-accelerated RAN, Sionna, fVDB, PhysX, Warp, NVblox |

# Accelerated Computing Is Sustainable Computing

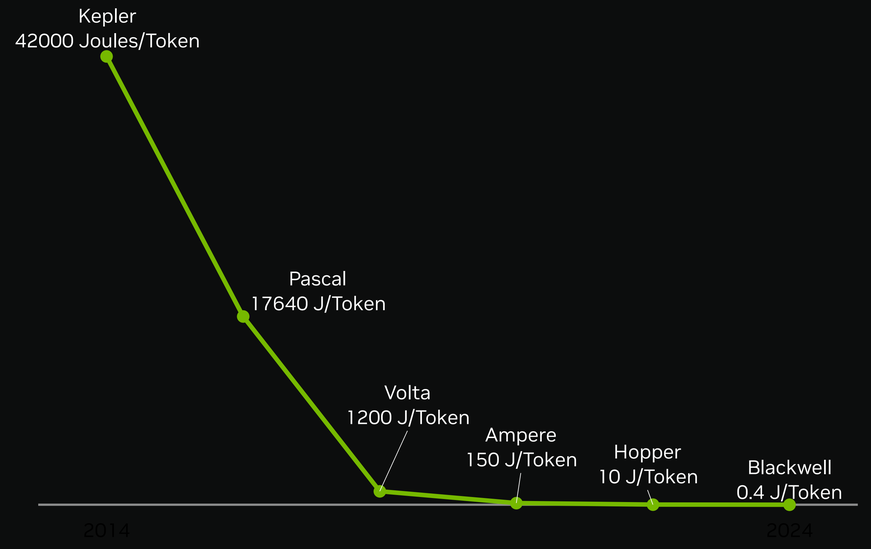
## Order of magnitude more energy efficient

Accelerated computing requires **higher peak power consumption** than CPUs, however, completes workloads **significantly faster and consumes less total energy**




Energy Usage in AI Inference

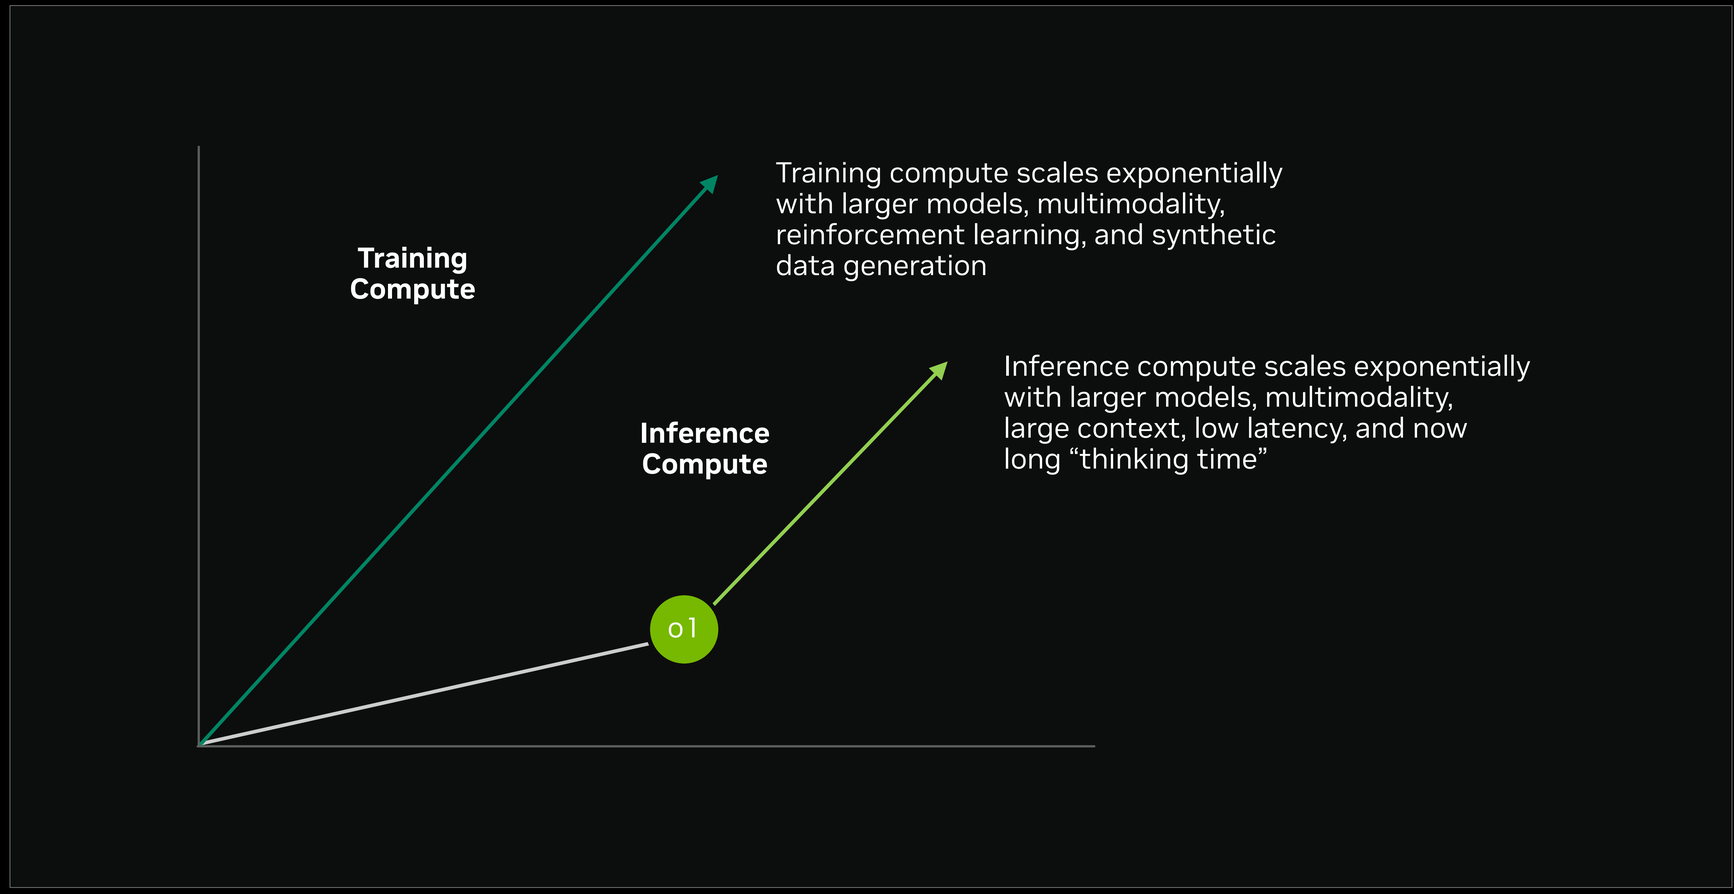
Accelerated computing enables **full-stack optimization** from algorithm to GPU architecture, such as Tensor Core Transformer Engine; **LLM energy efficiency improved 100,000X** in the past 10 years




GPT-MoE-1.8T energy per token

# AI Scaling Laws Drive Exponential Demand for Compute

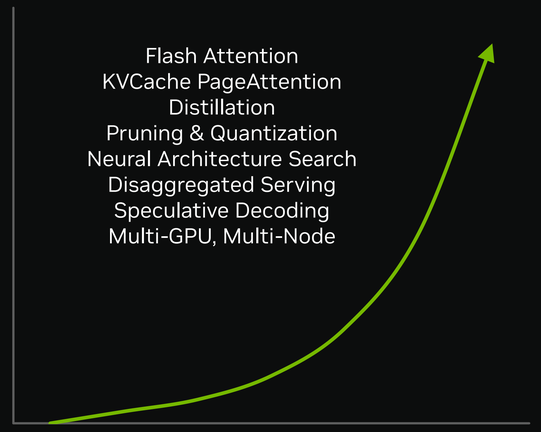
New OpenAI o1 Long “thinking time” creates a new way to scale



# NVIDIA Is the Leading Inference Platform

## Inference compute scales exponentially with “long thinking”

Hopper inference performance increased 5X in 1 year with rapid algorithm innovations enabled by rich NVIDIA CUDA ecosystem




Installed base & CUDA ➔ rapid software innovation ➔ performance ➔ lower inference cost ➔ increase demand ➔ increase installed base

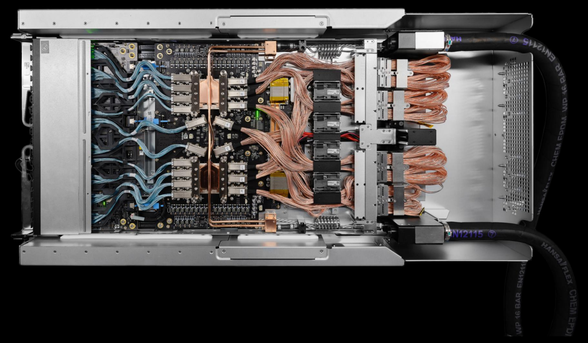


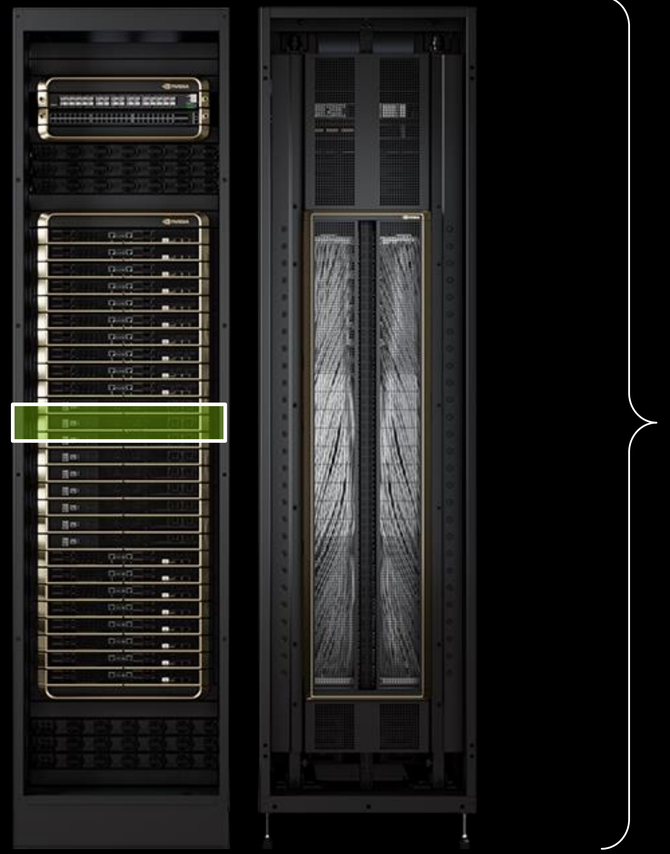
Inference compute scaling exponentially with large multimodal models, chain-of-thought, reasoning, agents, and low-latency responses



# NVIDIA NVLink Enables New Level of AI Training & Inference Scaling

NVLink Switch



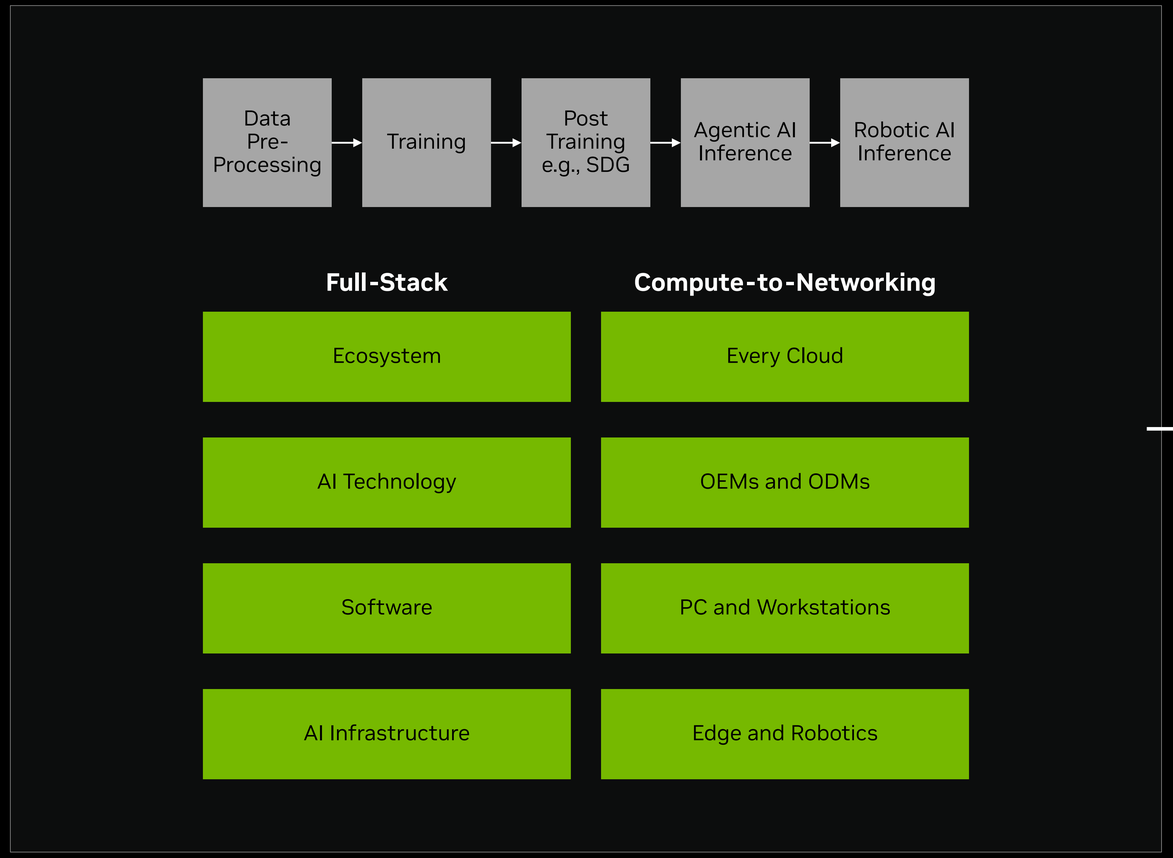


GB200 NVL72
NVLink Switches
130 TB/s All-to-All BW

One Giant Blackwell
1.44 EF FP4
576TB/s HBM3e

# NVIDIA AI Platform and Ecosystem Reaches Every Market

## Every workload to address the world's industries

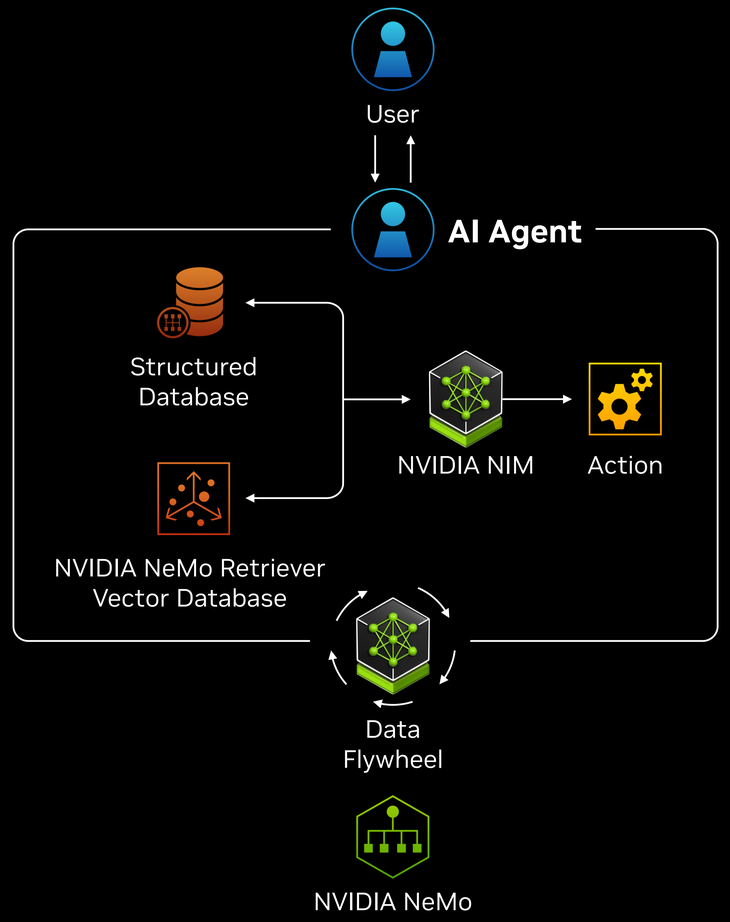



Accelerate Every Workload




Full-Stack, Entire AI Infrastructure

# NVIDIA AI Enterprise Enables IT Ecosystem With State-of-the-Art AI Models and Libraries to Build Agentic AI




## NVIDIA AI Enterprise Ecosystem

**System Integrators**

accenture Deloitte. . . .

**Enterprise ISVs**

amdocs cadence SAP servicenow SYNOPSYS . . .

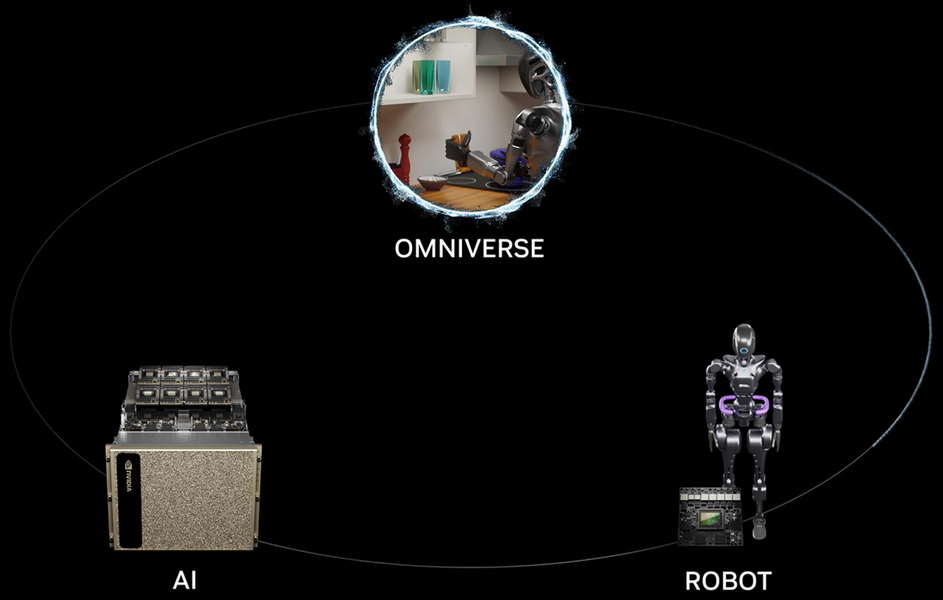
**Data Platforms**

CLOUDERA Dropbox NetApp ORACLE VAST . . .

**Cloud and On-Prem Infrastructure**

aws DELL Technologies Google Cloud Hewlett Packard Enterprise Microsoft Azure ORACLE CLOUD Infrastructure . . .

# NVIDIA Omniverse and AI Revolutionizing Manufacturing & Robotics




The next AI wave is physical AI—models that can perceive, understand, and interact with the physical world. Physical AI will embody robotic systems—from autonomous vehicles to industrial robots and humanoids, to warehouses and factories.

Three computers and software stacks are required to build physical AI:
NVIDIA AI on DGX to train the AI model, NVIDIA Omniverse on OVX to teach, test, and validate the AI model's skills, and NVIDIA AGX to run the AI software on the robot.

Enterprises license NVIDIA Omniverse at $4,500 per GPU per year.



100M Cars



Billions in Future



10M Factories



200K Warehouses

# Sovereign AI

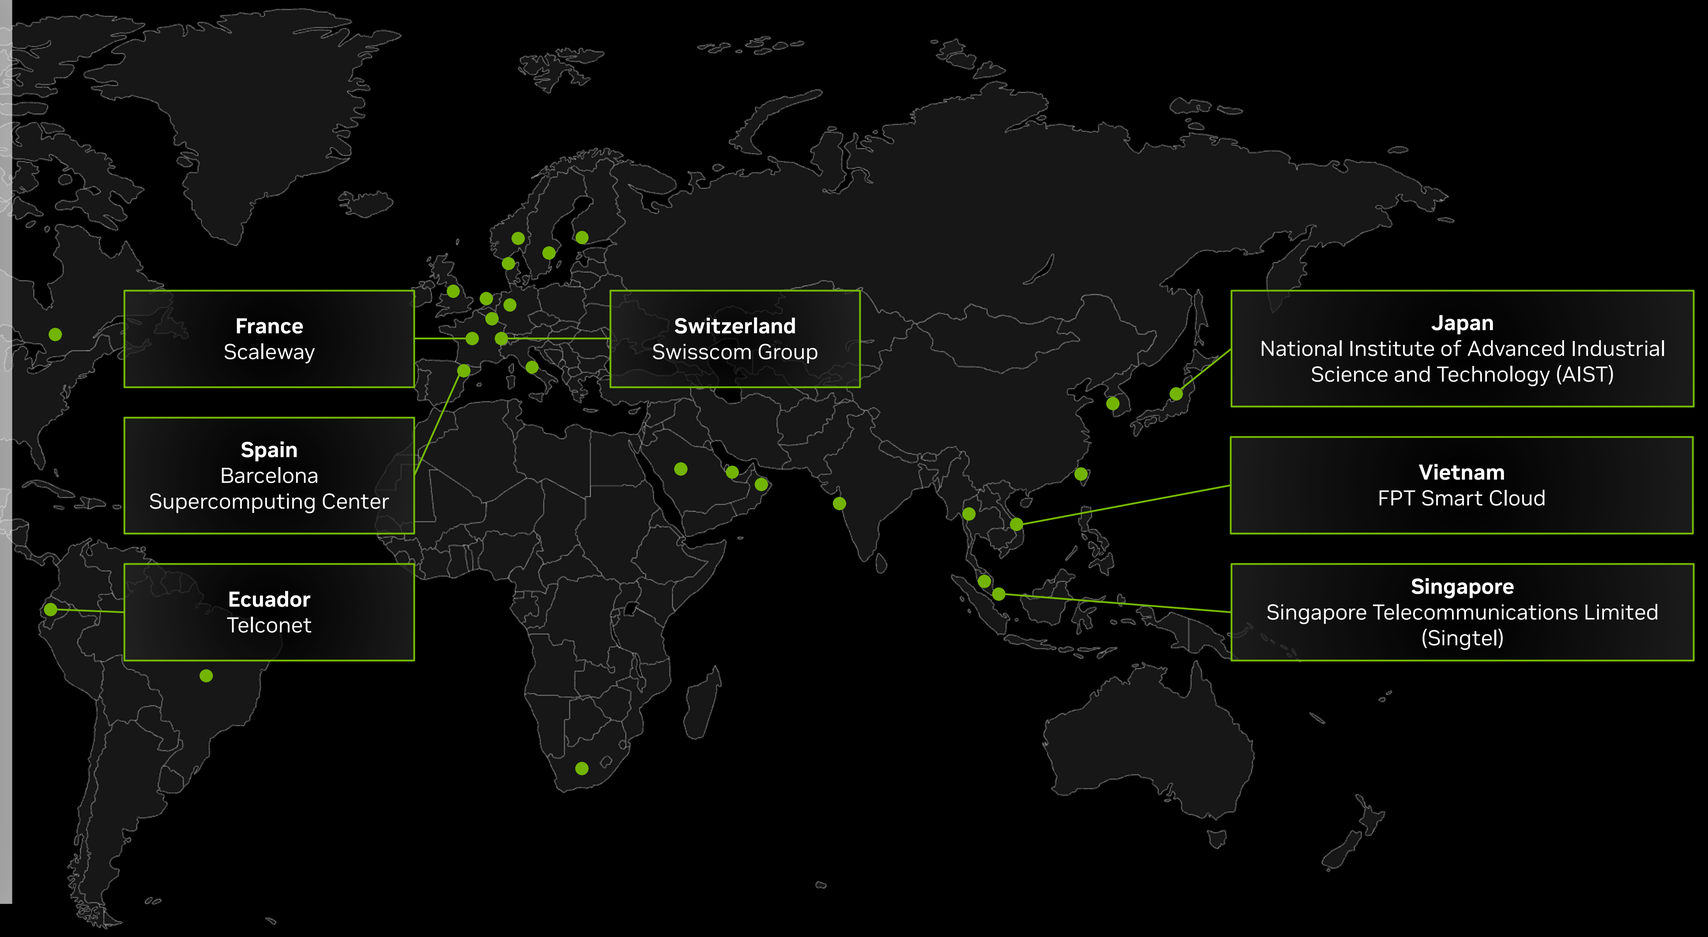
Nations produce AI using their own data, infrastructure, workforce, and business networks

## Sovereign AI

Nations are awakening to the imperative to produce AI using their own infrastructure, data, workforces, and business networks. Nations are building domestic computing capacity.

Some governments operate sovereign AI clouds in collaboration with state-owned telecommunications providers or utilities. Other governments partner with local cloud providers to deliver a shared AI computing infrastructure for public and private-sector use.

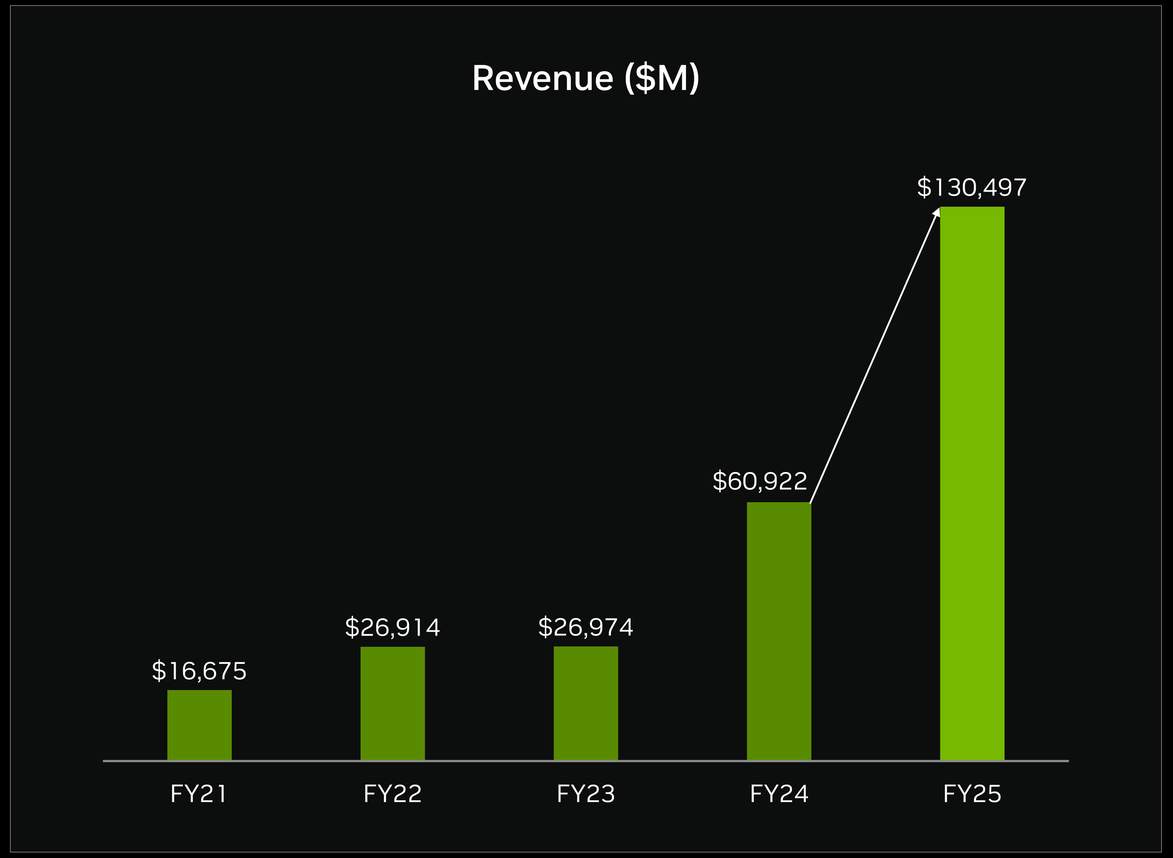
NVIDIA's ability to help build AI infrastructure with our end-to-end compute-to-networking technologies, full-stack software, AI expertise, and rich ecosystem of partners and customers allows sovereign AI and regional cloud providers to jump-start their countries' AI ambitions.

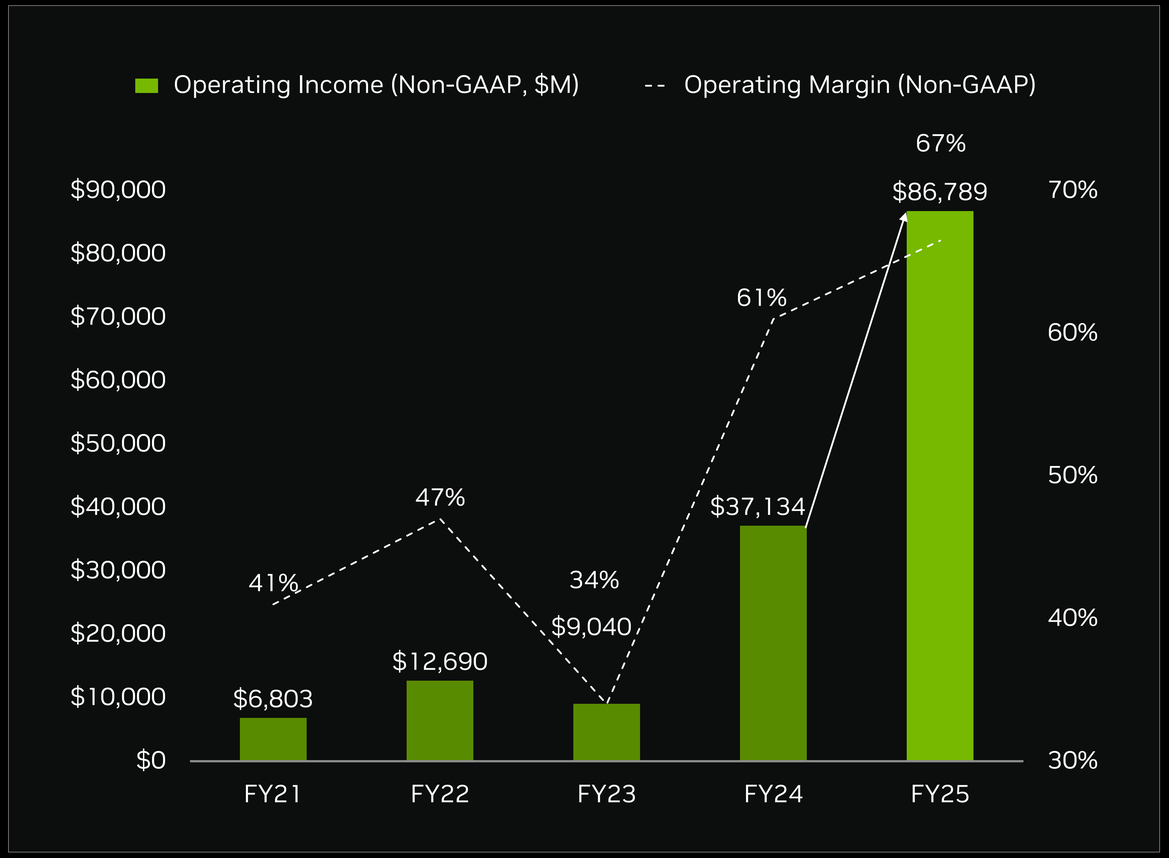



● Location of NVIDIA sovereign AI partners

NVIDIA

# Driving Strong and Profitable Growth

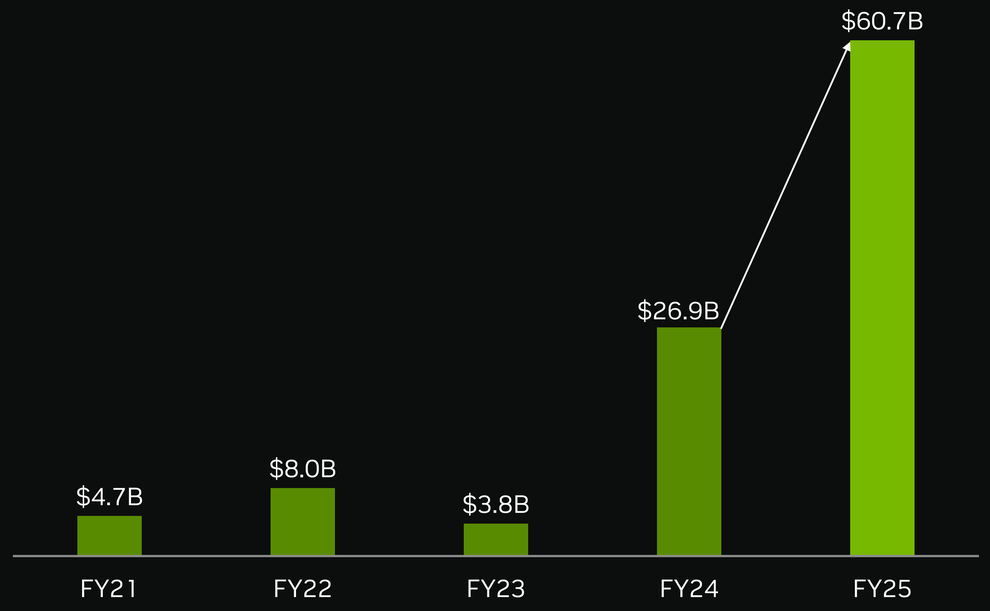



*Fiscal year ends in January. Refer to Appendix for reconciliation of Non-GAAP measures. Operating margins rounded to the nearest percent.*

# Strong Cash Flow Generation

## Free Cash Flow (Non-GAAP)




*Fiscal year ends in January. Refer to Appendix for reconciliation of Non-GAAP measures.*

## Capital Allocation

### Share Repurchase

Utilized $33.7B of cash for repurchases in FY25

$38.7B remaining authorization as of the end of Q4

### Dividend

$834M in FY25

Dividend increased by 150% in Q2 FY25

Plan to Maintain[1]

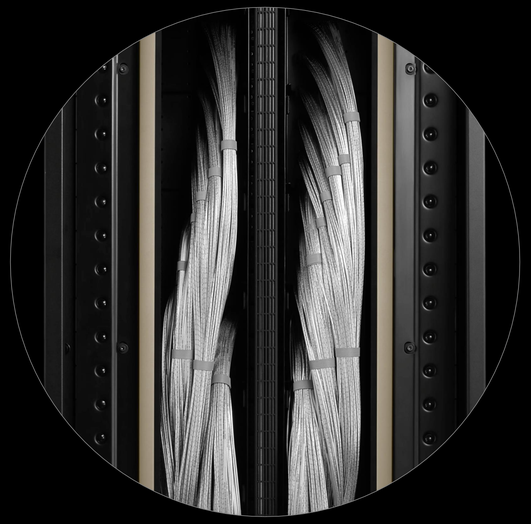
### Strategic Investments

Growing Our Talent

Platform Reach and Ecosystem

[1] *Subject to continuing determination by our Board of Directors.*

# Our Market Platforms at a Glance



## Data Center

88% of FY25 Revenue

**FY25 Revenue $115.2B**
5-YR CAGR 108%

DGX/HGX/MGX/IGX systems

GPU | CPU | DPU | Networking
NVIDIA AI software



## Gaming

9% of FY25 Revenue

**FY25 Revenue $11.4B**
5-YR CAGR 16%

GeForce GPUs for PC gaming

GeForce NOW cloud gaming



## Professional Visualization

1% of FY25 Revenue

**FY25 Revenue $1.9B**
5-YR CAGR 9%

NVIDIA RTX GPUs
for workstations

Omniverse software




## Automotive

1% of FY25 Revenue

**FY25 Revenue $1.7B**
5-YR CAGR 19%

DRIVE Hyperion sensor architecture
with AGX compute

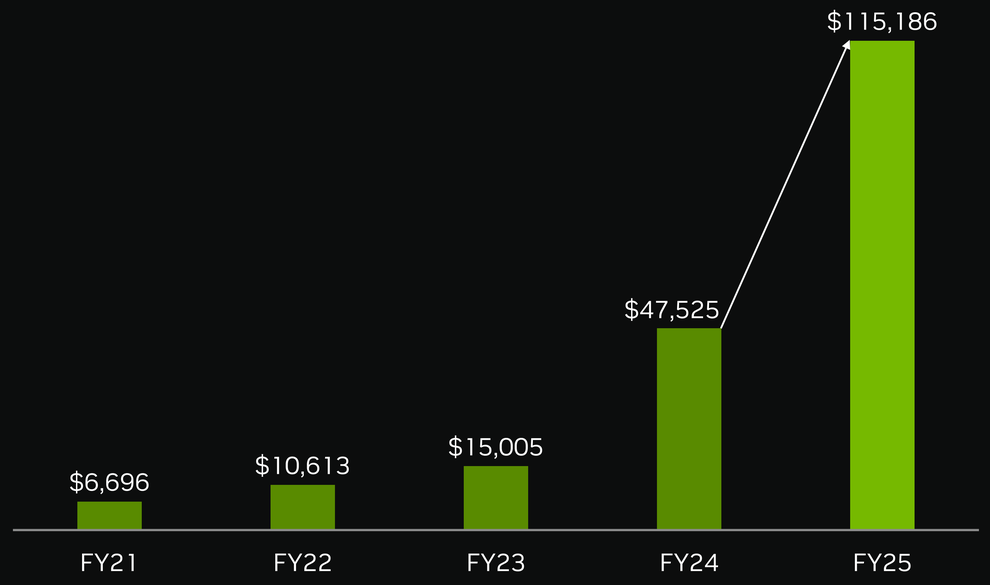
DRIVE AV & IX full-stack software
for ADAS, AV, and AI cockpit

# Data Center

## The leading accelerated computing platform

### Revenue ($M)

108% 5-YR CAGR
Through FY25




### Leader in AI and HPC

No. 1 in AI training and inference

Used by all hyperscalers, major cloud computing providers, and over 40,000 companies

Powers over 75% of the TOP500 supercomputers

### Growth Drivers

Broad data center platform transition from general-purpose to accelerated computing

Emergence of AI factories—optimized for refining data and training, inferencing, and generating AI

Broader and faster product launch cadence to meet a growing and diverse set of AI opportunities

NVIDIA AI Enterprise/NIM for building and running enterprise AI applications

# One-Year Rhythm | Supercluster Scale | Full-Stack | CUDA Everywhere

## Supercharge AI scaling law

X-Factors

X-Factors

X-Factors

### Hopper

Hopper GPU
6S HBM3

Hopper+ GPU
6S HBM3e

Grace CPU

NVLink Switch
900 GB/sec

CX7
SuperNIC

BF3
SuperNIC

Quantum-X400
Infiniband Switch

### Blackwell




### Blackwell-Ultra

Blackwell Ultra GPU
288GB HBM3e
More AI FLOPS

CX8
SuperNIC

Spectrum Ultra X800
Ethernet Switch 512-Radix

### Rubin

Rubin GPU
8S HBM4

Rubin Ultra GPU
12S HBM4

Vera CPU

NVLink 6 Switch
3600 GB/sec

CX9 SuperNIC
1600 Gb/sec

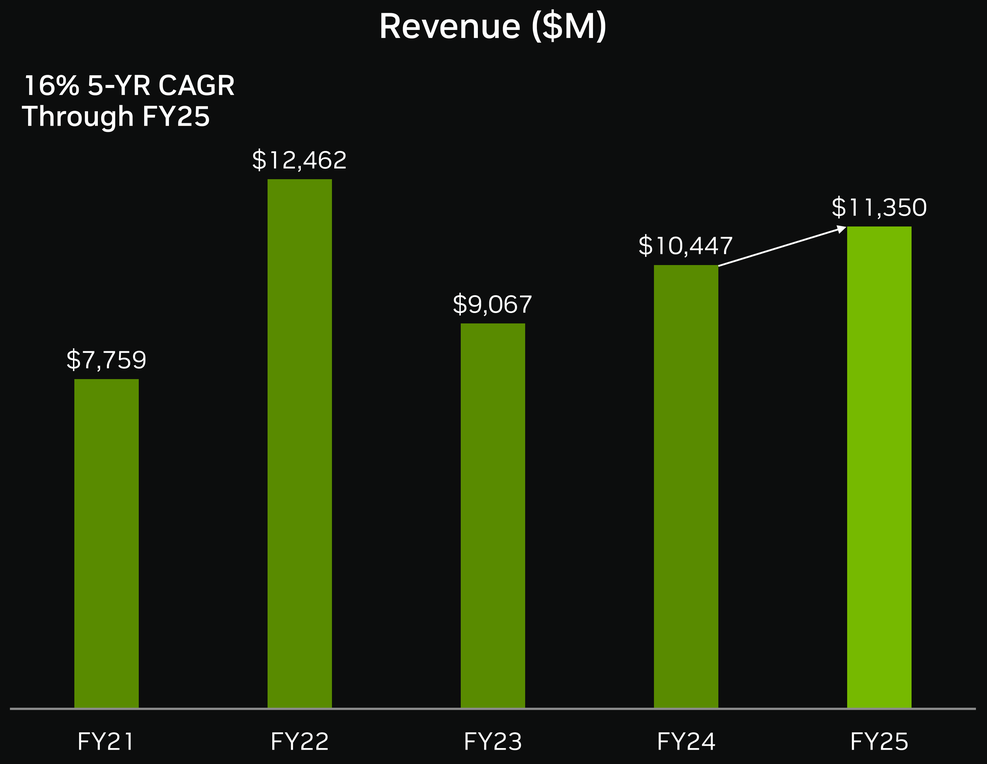
X1600
IB/Ethernet Switch

2022 2023 2024 2025 2026 2027

NVIDIA

# Gaming

## GeForce—world's largest gaming platform




### Leader in PC Gaming

Strong No. 1 market position

15 of the top 15 most popular GPUs on Steam

Leading performance and innovation

200M+ gamers on GeForce

### Growth Drivers

Rising adoption of NVIDIA RTX in games

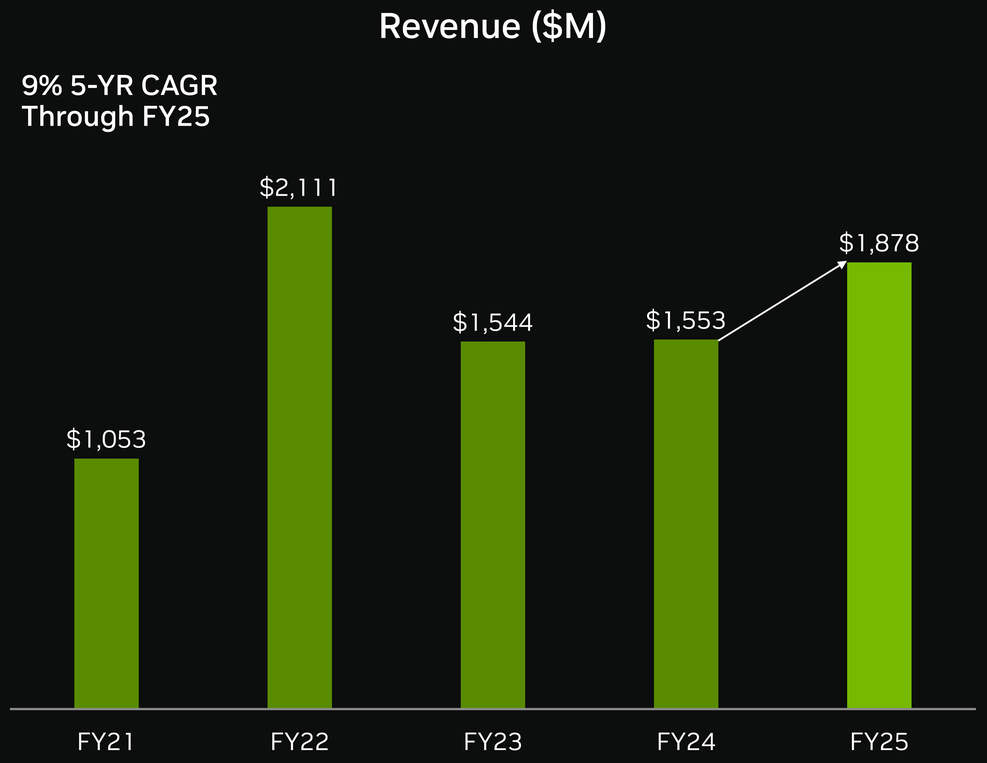
Expanding universe of gamers and creators

Gaming laptops and generative AI on PCs

GeForce NOW cloud gaming

# Professional Visualization

## Workstation graphics




### Leader in Workstation Graphics

95%+ market share in graphics for workstations

45M designers and creators

Strong software ecosystem with over 100 RTX accelerated and supported applications

### Growth Drivers

Generative AI adoption across design and creative industries

Enterprise AI development, model fine-tuning, cross-industry

Ray tracing revolutionizing design and content creation

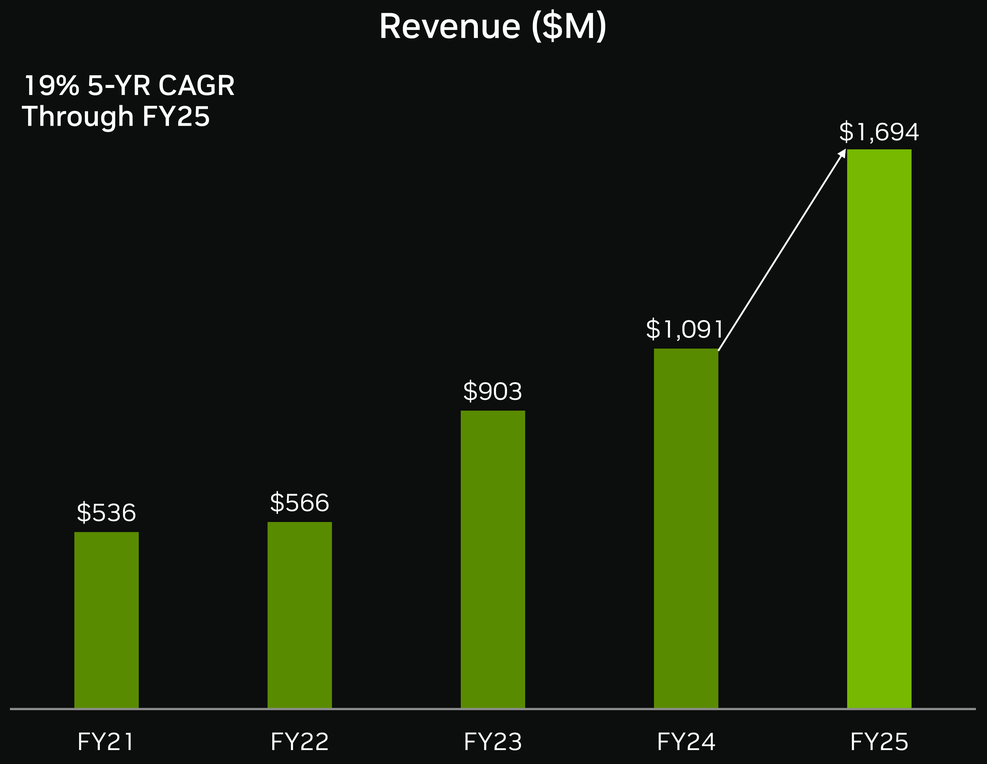
Expanding universe of designers and creators

Omniverse for digital twins and collaborative 3D design

Hybrid work environments

# Automotive

## Autonomous vehicles and AI cockpits




### Leader in Autonomous Driving

NVIDIA DRIVE an end-to-end autonomous vehicle (AV) and AI cockpit platform featuring a full software stack and powered by NVIDIA SoCs (systems-on-a-chip) in vehicles

DRIVE Orin SoC ramp began in FY23

Next-generation DRIVE Thor SoC ramp to begin in FY26

Over 40 customers including 20 of top 30 EV makers, 7 of top 10 truck makers, 8 of top 10 robotaxi makers

### Growth Drivers

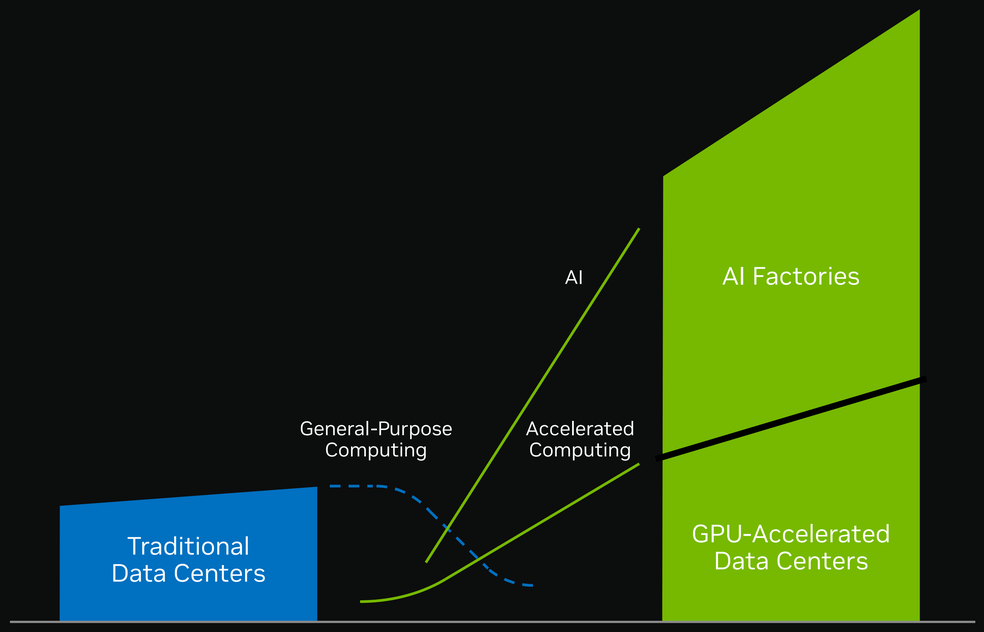
Adoption of centralized car computing and software-defined vehicle architectures

AV software and services:

Mercedes-Benz

Jaguar Land Rover

# Accelerated Computing and Generative AI Create Trillion-Dollar Opportunities



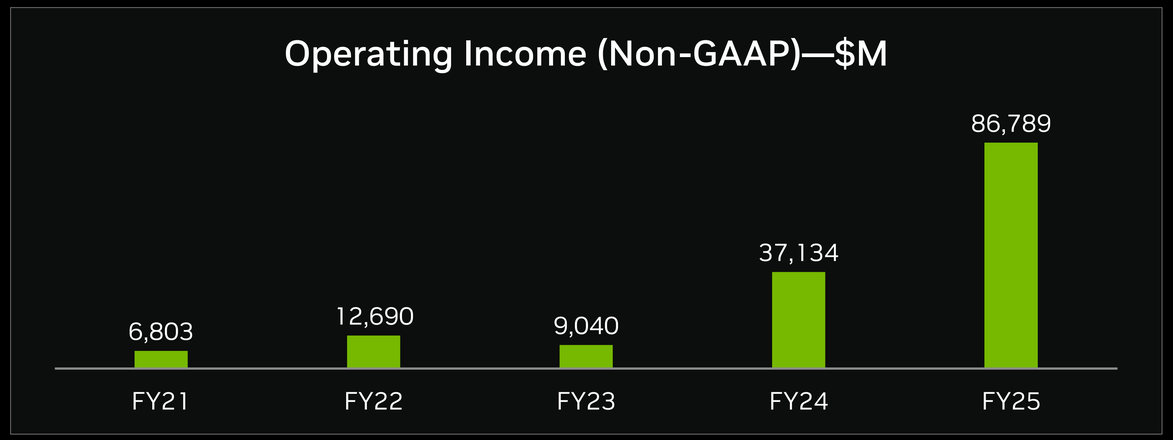

The $1T installed base of general-purpose CPU data center infrastructure is being modernized to a new GPU-accelerated computing paradigm.

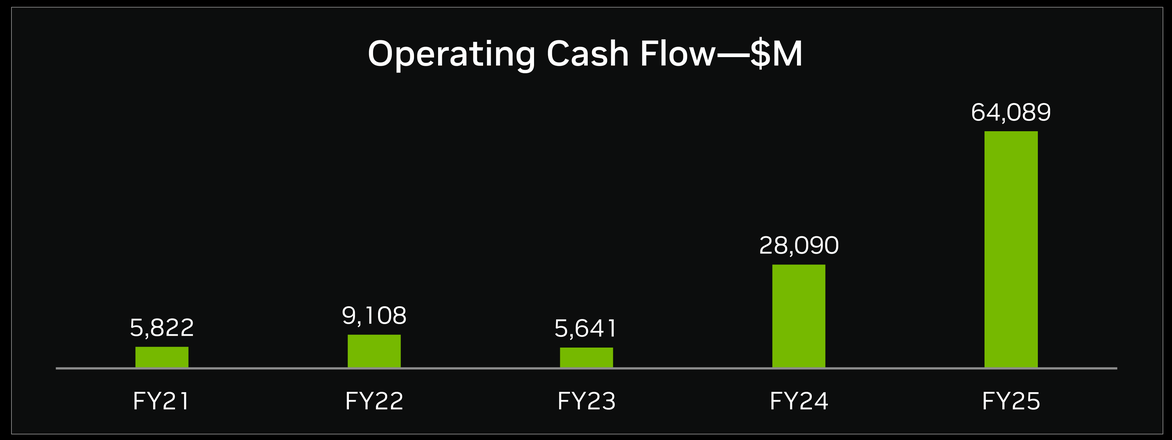
The entire computing stack has been reinvented—from CPU to GPU, from coding to machine learning, from software to generative AI. Computers generate intelligence tokens, a new commodity.

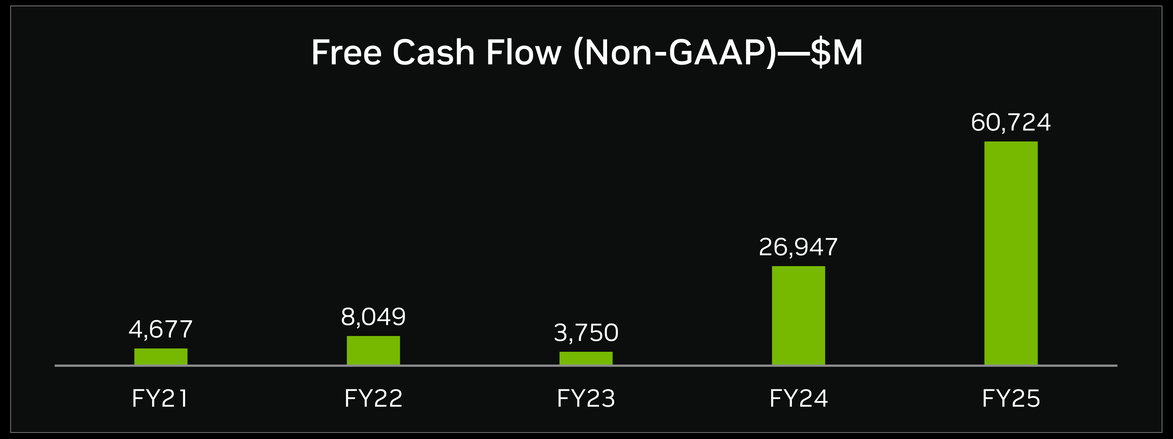
A new type of data center, AI factories, is expanding the data center footprint to $2T and beyond in the coming years. Eventually, companies in every industry will operate AI factories as the digital twin of their workforce, manufacturing plants, and products. A new industrial revolution has begun.

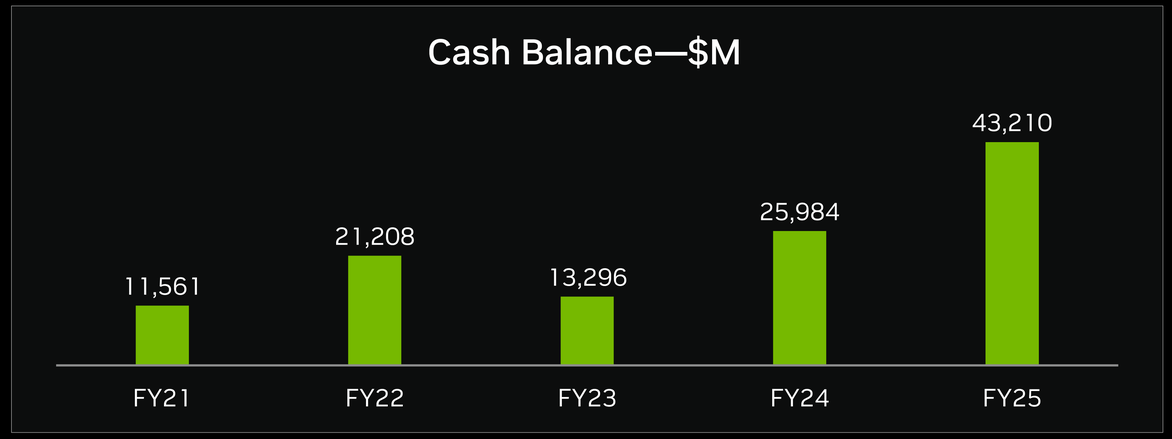
# Financials

# Annual Cash & Cash Flow Metrics




*Cash balance is defined as cash and cash equivalents plus marketable securities*
*Refer to Appendix for reconciliation of non-GAAP measures*

# Corporate Sustainability

## Environmentally Conscious



NVIDIA Blackwell GPUs are as much as 20X more energy efficient than CPUs for certain AI and HPC workloads



On track to source 100% renewable electricity for offices and data centers under operational control by end of FY25



On track to engage manufacturing suppliers comprising at least 67% of scope 3 category 1 GHG emissions with the goal of effecting supplier adoption of science-based targets by end of FY26

## A Place for People to Do Their Life's Work

"Best Places to Work"
**GLASSDOOR**

"America's Most Sustainable Companies"
**BARRON'S**

"America's 100 Best Companies to Work For"
**FORTUNE**

"America's Most Responsible Companies"
**NEWSWEEK**

## Management

Fast Company Magazine's World's Most Innovative Companies

Fortune's World's Most Admired Companies

Time Magazine's 100 Most Influential Companies

Wall Street Journal's Management Top 250

## Corporate Governance

92% of directors are independent

# Reconciliation of Non-GAAP to GAAP Financial Measures

# Reconciliation of Non-GAAP to GAAP Financial Measures

| Operating Income and Margin ($ in Millions and Margin Percentage) | Non-GAAP | Acquisition Termination Cost | Acquisition-Related and Other Costs (A) | Stock-Based Compensation (B) | Other (C) | GAAP |
|---|---|---|---|---|---|---|
| FY 2021 | $6,803 | — | (836) | (1,397) | (38) | $4,532 |
| | 40.8% | — | (5.0) | (8.4) | (0.2) | 27.2% |
| FY 2022 | $12,690 | — | (636) | (2,004) | (9) | $10,041 |
| | 47.2% | — | (2.5) | (7.4) | — | 37.3% |
| FY 2023 | $9,040 | (1,353) | (674) | (2,710) | (79) | $4,224 |
| | 33.5% | (5.0) | (2.5) | (10.0) | (0.3) | 15.7% |
| FY 2024 | $37,134 | — | (583) | (3,549) | (30) | $32,972 |
| | 61.0% | — | (1.0) | (5.8) | (0.1) | 54.1% |
| FY 2025 | $86,789 | — | (602) | (4,737) | 3 | $81,453 |
| | 66.5% | — | (0.5) | (3.6) | — | 62.4% |

*A. Consists of amortization of acquisition-related intangible assets, inventory step-up, transaction costs, compensation charges, and other costs*
*B. Stock-based compensation charge was allocated to cost of goods sold, research and development expense, and sales, general and administrative expense*
*C. Comprises of legal settlement cost, contributions, restructuring costs and assets held for sale related adjustments*

# Reconciliation of Non-GAAP to GAAP Financial Measures (contd.)

| ($ in Millions) | Free Cash Flow | Purchases Related to Property and Equipment and Intangible Assets | Principal Payments on Property and Equipment and Intangible Assets | Net Cash Provided by Operating Activities |
|---|---|---|---|---|
| **FY 2021** | $4,677 | 1,128 | 17 | $5,822 |
| **FY 2022** | $8,049 | 976 | 83 | $9,108 |
| **FY 2023** | $3,750 | 1,833 | 58 | $5,641 |
| **FY 2024** | $26,947 | 1,069 | 74 | $28,090 |
| **FY 2025** | $60,724 | 3,236 | 129 | $64,089 |


NVIDIA